

# Lecture 3 - MapReduce Details

---

**BDAT 1002**

Review

## Q1

---

- Which of the below is not a file system?
  - HDFS
  - ext4
  - FS4
  - ext3

## Q2

---

- Assume the cluster (block) size of NTFS is 32 KB. If you are trying to store a file of size 5 KB in your file system, how much space will that file occupy in your hard disk?
  - 5 KB
  - 32 KB
  - 64 KB
  - 64 MB

## Q3

---

- Assume the block size of HDFS is 64 MB and the block size of underlying file system is 1 MB. If you are trying to store a file of size 5 MB in your file system, how much space will that file occupy in your hard disk?
  - 64 MB
  - 5 MB
  - 128 MB
  - 300 MB

## Q4

---

- The concept of HDFS is based on which of the below file system?
  - NTFS
  - ext4
  - ext3
  - GFS

## Q5

---

- What is the default replication factor in HDFS?
  - 5
  - 4
  - 3
  - 2

## Q6

---

- What is the maximum RAM that a 32 bit machine can support?



# MapReduce Introduction

# MapReduce Introduction

---

- Example problem: Governor of California wants you to do a census to find the population of California by city
  - You have all the resources you want
- What is your process?
  - Can you do it yourself?
- Sensible thing to do is "divide and conquer"

# MapReduce Introduction

---

- For simplicity, let's say there are only three cities
- Have a person in charge of each of the cities
- Each person responsible for finding the population of each city

SFO



SJOSE



LA



# MapReduce Introduction

- Each person goes home by home and write down the number of people in the household

SFO



SFO 5  
SFO 3  
SFO 4  
SFO 1

SJOSE



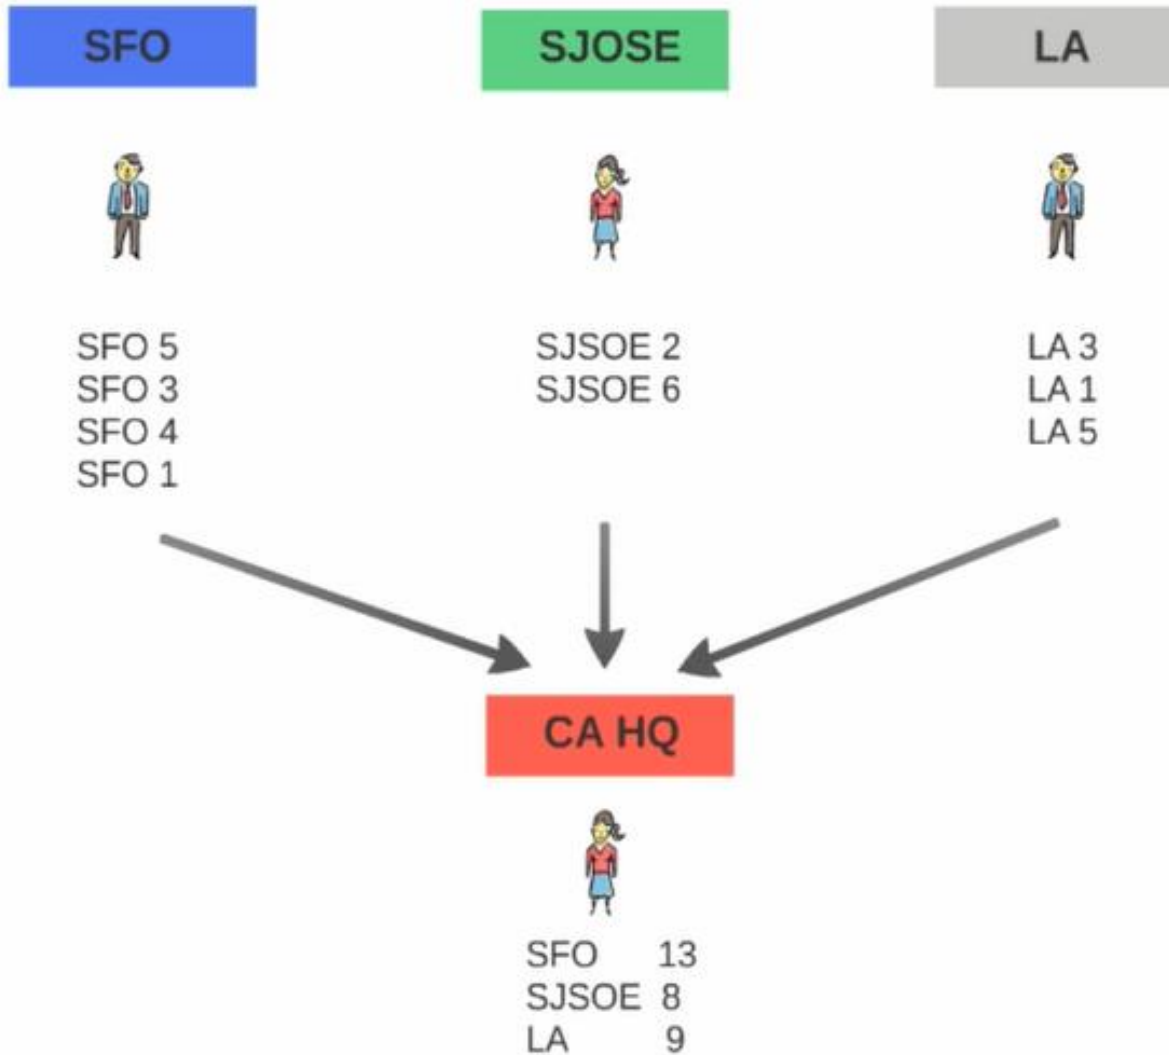
SJSOE 2  
SJSOE 6

LA



LA 3  
LA 1  
LA 5

# MapReduce Introduction



# MapReduce Introduction

---


- Next year, you are now asked to do the same thing **BUT** with **half the time**
- What to do?
- Remember you have all the resources you want?

# MapReduce Introduction

- Double the number of people doing count for each city



# MapReduce Introduction





# MapReduce Introduction

---

- There is one crucial detail here
- You want your census takers for each city (ie SFO 1 and SFO2) send the results to **either HQ1 or HQ2**
- We don't want results spread across two headquarters
  - Remember we wanted population by city
- So there must be instructions for census takers in one city to coordinate and send data to same headquarters

# MapReduce Introduction

---

- We can expand this model, more resources, less time
- This model works and it can scale
- The model is called MapReduce

# What is MapReduce?

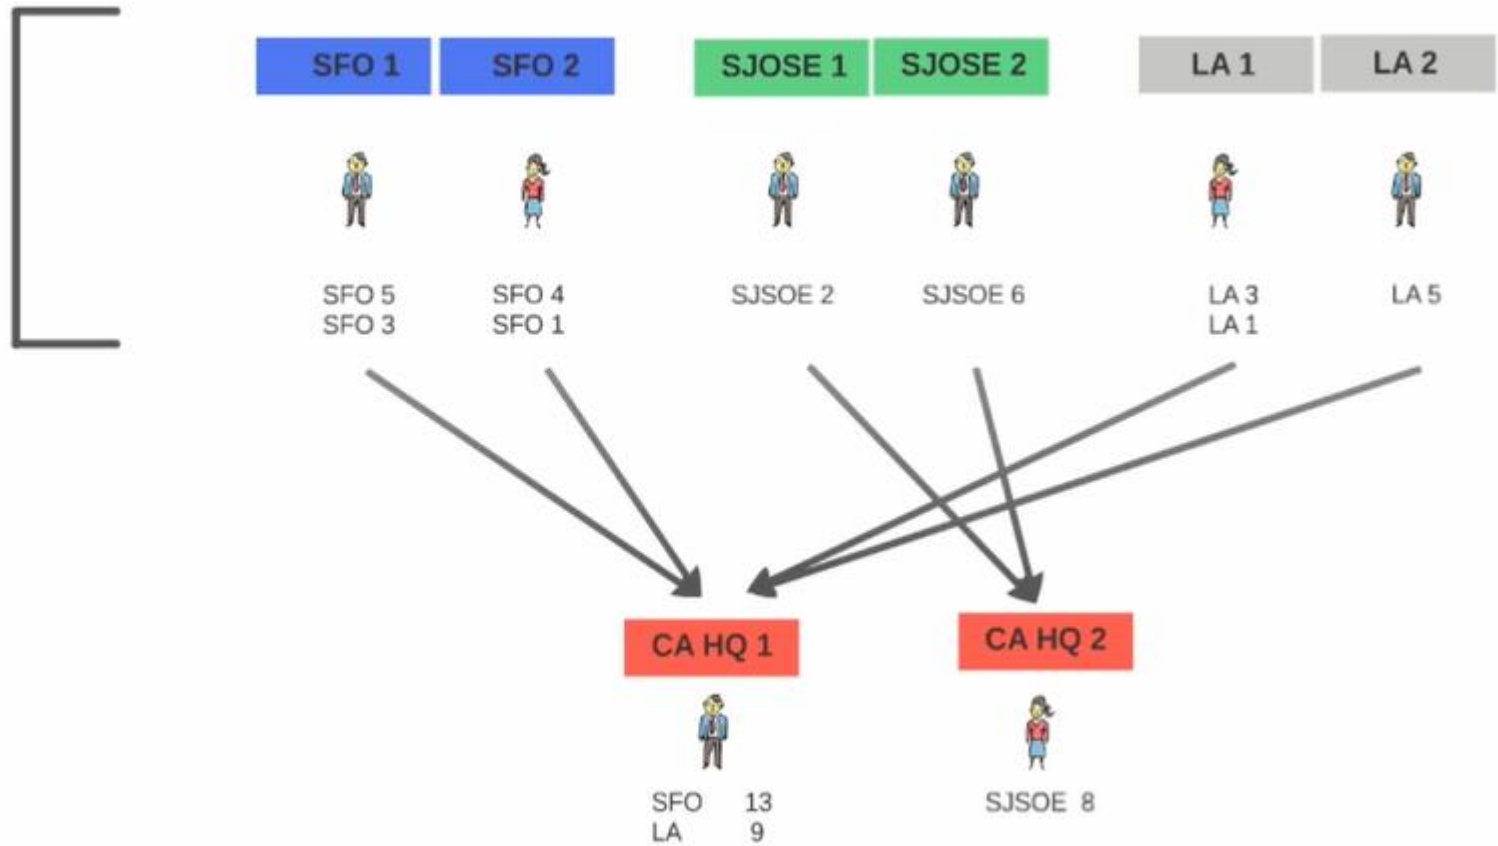
---

- MapReduce is a programming model for ***distributive computing***
  - It is not a programming language

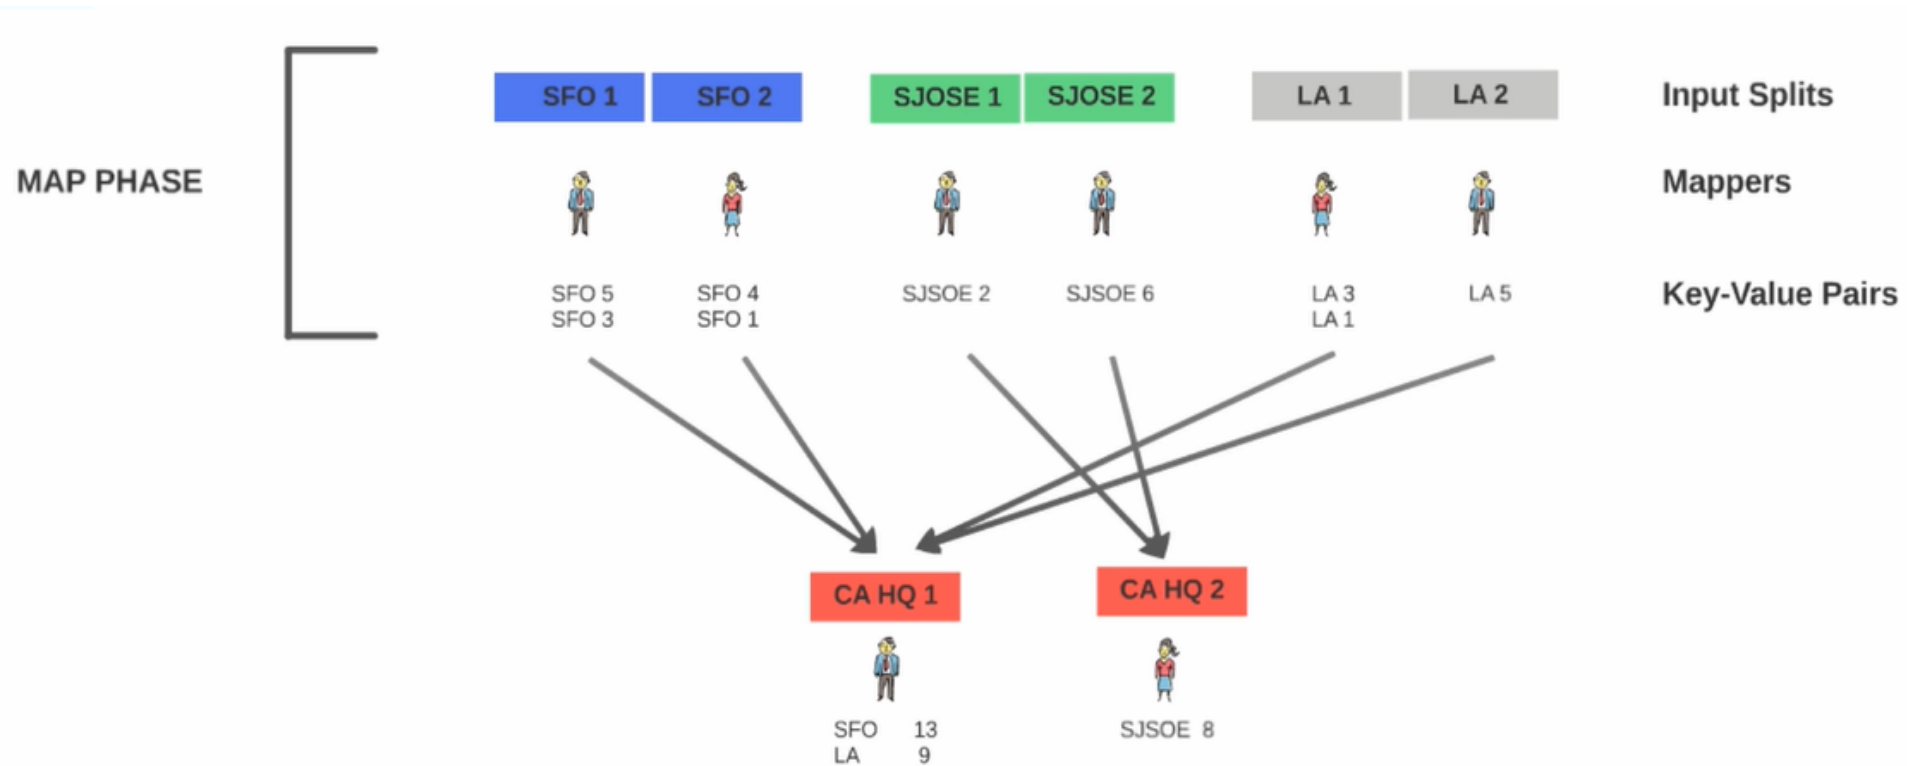
# Phases in MapReduce

# Phases of MapReduce

MAP PHASE



# Phases of MapReduce

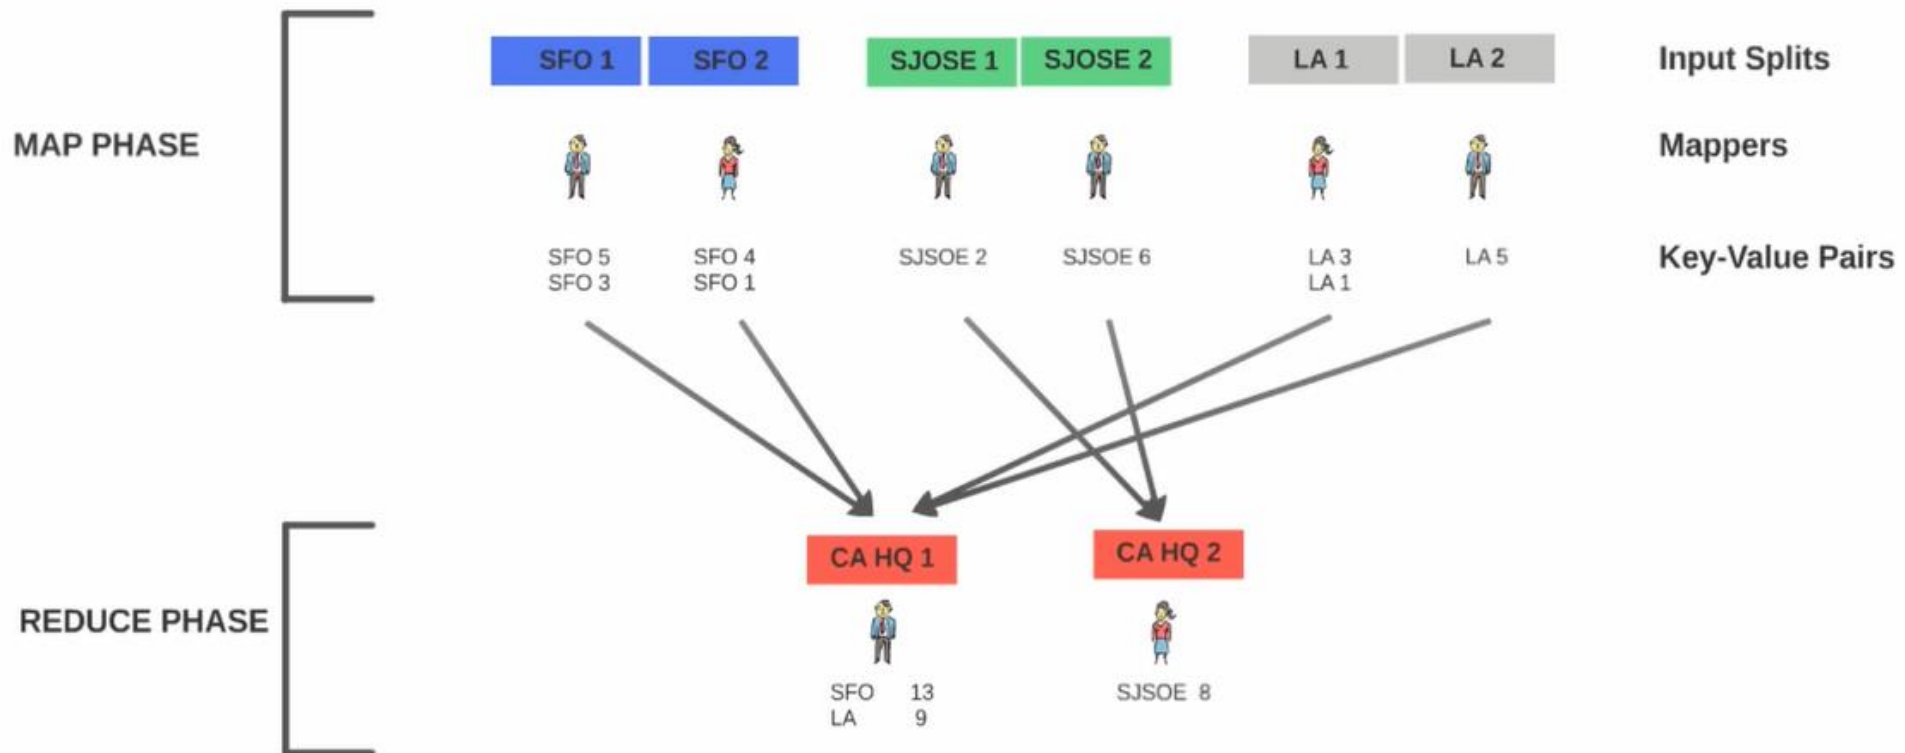


## Phases of MapReduce

---

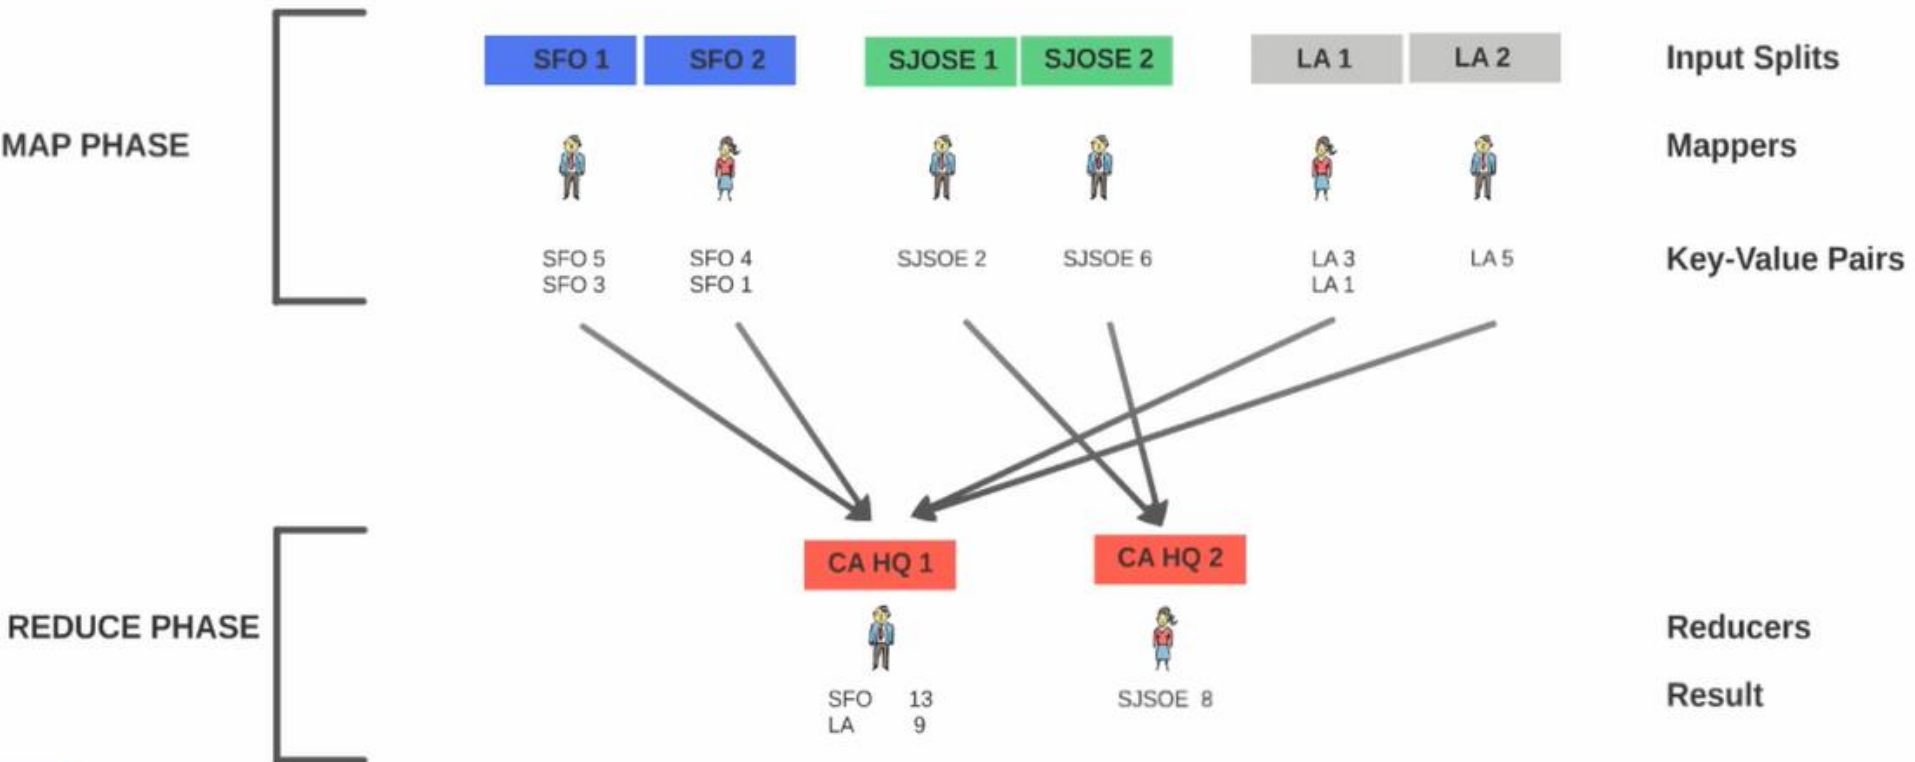
- The phase where individuals collect the population of individual cities is called the **Map Phase**
- Individual person involved in the calculations is called **the Mapper**
- The city or part of the city they are working in is called **Input splits**
- The output from each pair is **key-value pair**

# Phases of MapReduce

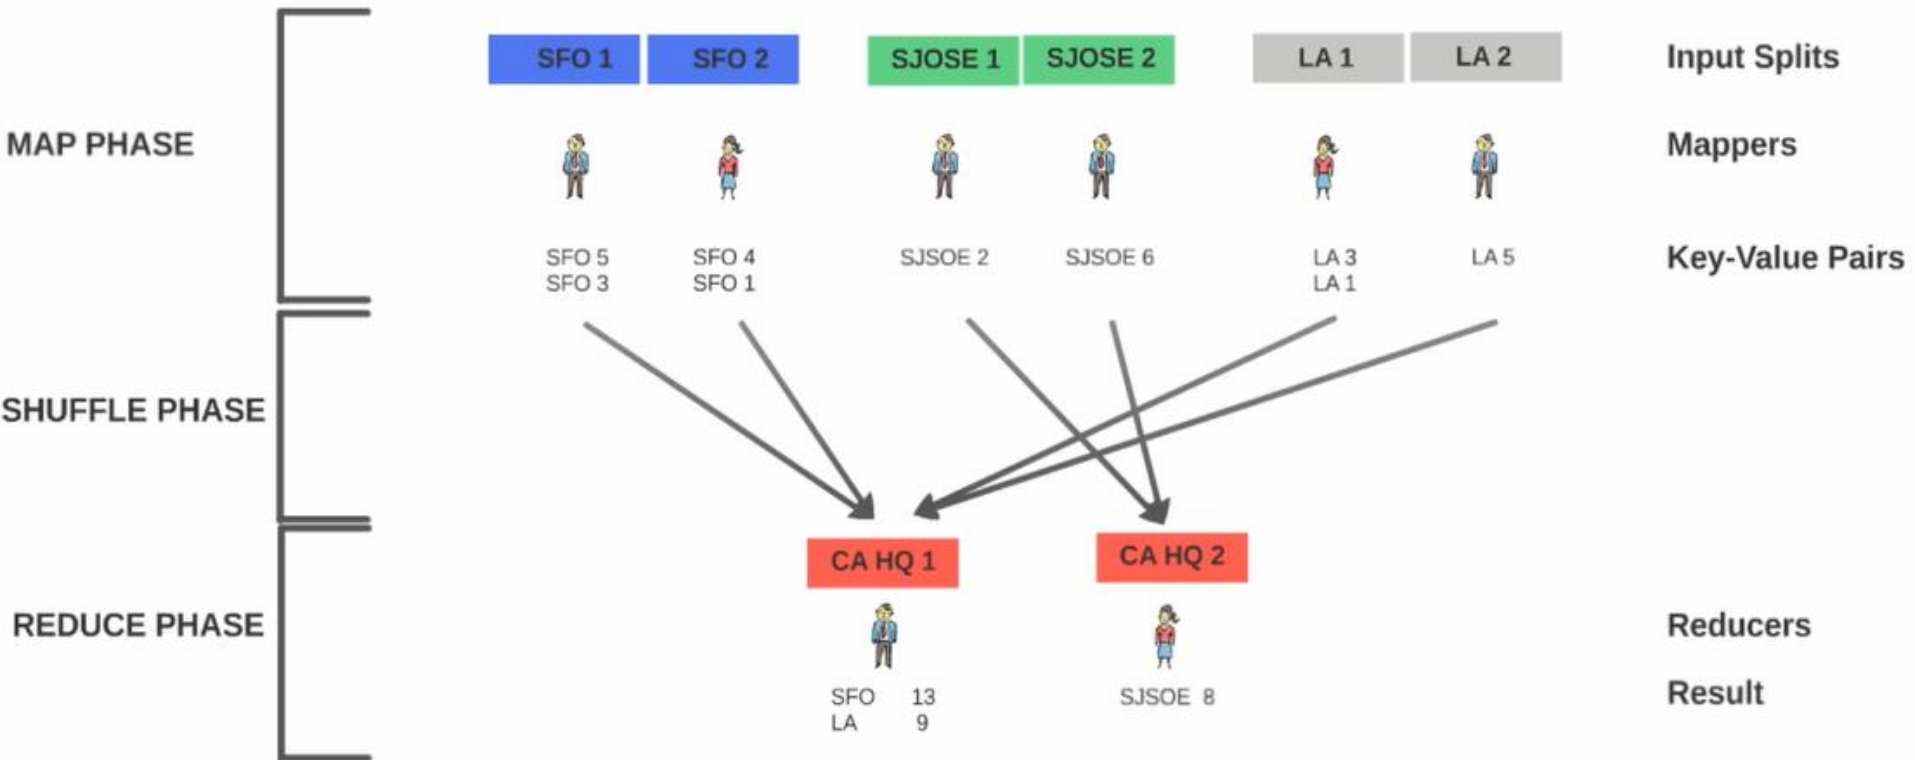




# Phases of MapReduce



# Phases of MapReduce



# Phases of MapReduce

---

- The phase you aggregate the results for each city is called the **reduce phase**
- The individuals working in the headquarters are known as the **reducers**
- Each reducer will produce a **result set**
- The phase where the values for the different mappers is copied or transferred to the reducers is known as the **shuffle phase**

# Summary

---

- MapReduce is a programming model for processing large data
  - Distributed processing
  - Developed at Google
  - Adapted by Hadoop
- Can be implemented in any programming language
  - Hadoop supports many languages
  - Java, C++, Python, Scala

# MapReduce in Detail

# Sample Big Data Problem

---

- Sample stock dataset
- Each record has symbol, date, open, close, max, min, volume, etc
- Want to find the maximum closing price for each symbol

# Sample Big Data Problem

---

- Each line in the data set is a **record**

```
ABCSE,B7J,2008-10-28,6.48,6.74,6.22,6.72,44300,5.79
ABCSE,B7J,2008-10-27,6.21,6.78,6.21,6.40,55200,5.51
ABCSE,B7J,2008-10-24,6.39,6.66,6.21,6.40,67400,5.51
ABCSE,B7J,2008-10-23,6.95,6.95,6.50,6.59,59400,5.68
ABCSE,B7J,2008-10-22,6.92,7.17,6.80,6.80,55300,5.86
ABCSE,B7J,2008-10-21,7.20,7.30,7.10,7.10,54400,6.11
ABCSE,B7J,2008-10-20,6.94,7.31,6.94,7.12,45700,6.13
ABCSE,B7J,2008-10-17,6.43,6.93,6.42,6.90,57700,5.94
ABCSE,B7J,2008-10-16,6.61,6.69,6.21,6.53,83200,5.62
ABCSE,B7J,2008-10-15,6.84,6.90,6.36,6.36,78900,5.48
ABCSE,B7J,2008-10-14,7.15,7.32,6.93,6.96,74700,5.99
ABCSE,B7J,2008-10-13,6.00,6.57,6.00,6.57,75700,5.66
ABCSE,B7J,2008-10-10,5.05,5.72,4.79,5.72,158400,4.93
ABCSE,B7J,2008-10-09,6.30,6.41,6.00,6.02,140500,5.18
ABCSE,B7J,2008-10-08,5.60,6.47,5.60,6.28,292000,5.41
ABCSE,B7J,2008-10-07,7.59,7.59,6.66,6.69,89900,5.76
ABCSE,B7J,2008-10-06,7.83,7.90,7.00,7.40,159600,6.37
```

# Sample Big Data Problem

---

- How to solve this problem?
  - Forget about MapReduce etc
- Parse each record
- Check closing price of stock
- If price higher than a current maximum price, change price
- Else go to next record
- Stop when you reach last line (record)



# Sample Big Data Problem

---

- As we discussed in Lecture 1, the problem with this approach is time
- With no parallelization, it may take very long time to complete task

## Solution in MapReduce World

---

- We will go through the phases we discussed in previous section
  - Kind of abstract with census example
  - Now we want to see details with an actual problem
- First talk about **Map phase**

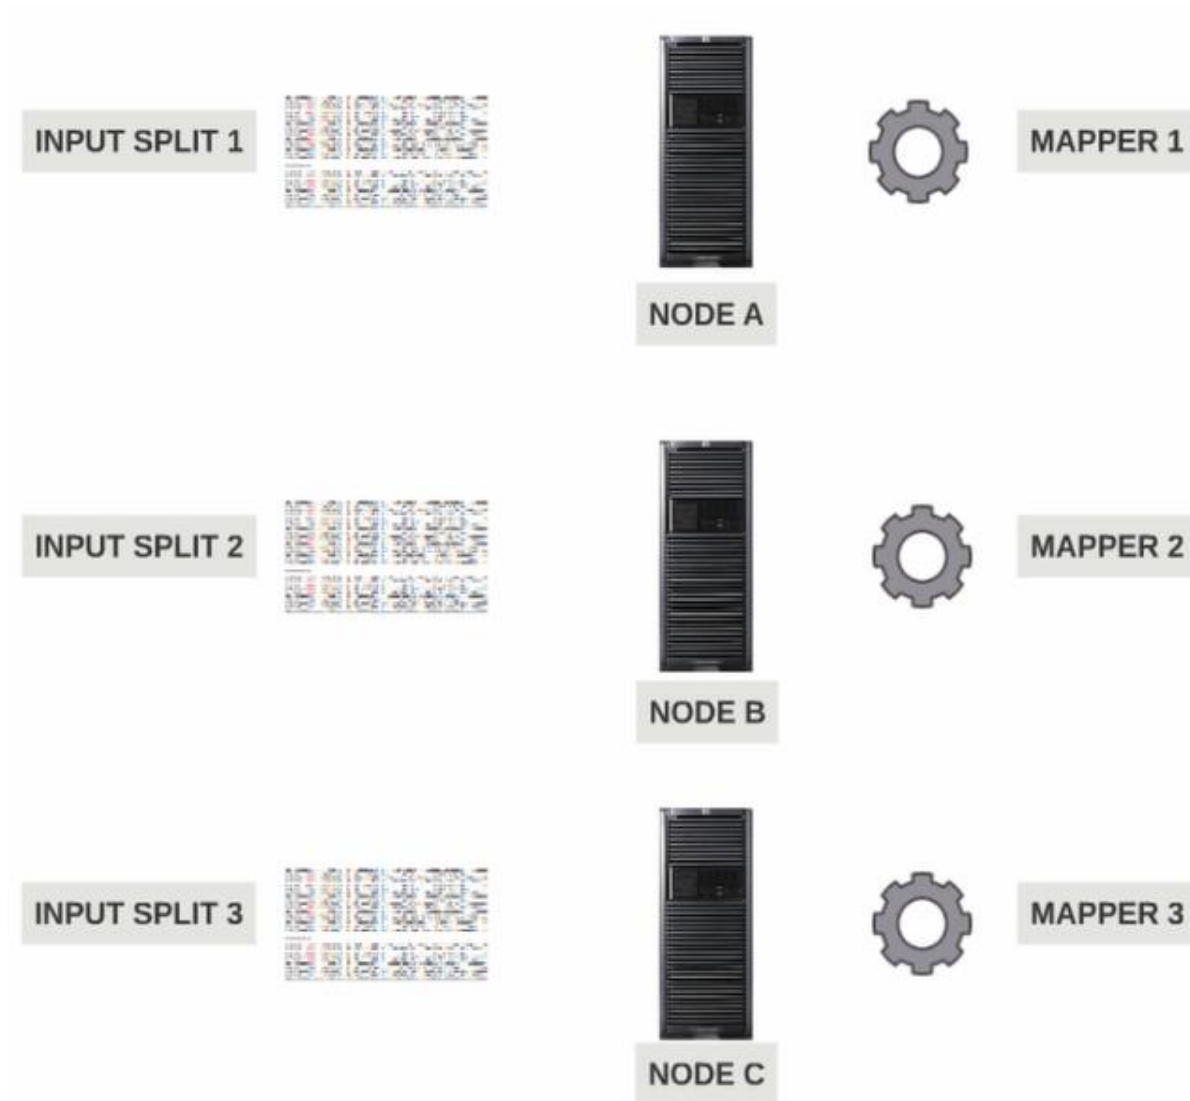
## Map Phase

---

- Divide the data set into chunks
- Have a separate process working on each chunk of the data
- Some technical jargon now
  - Chunks are called **input splits**
  - The process working on the chunks are called **mappers**

# Graphic View

---



## Map Phase

---

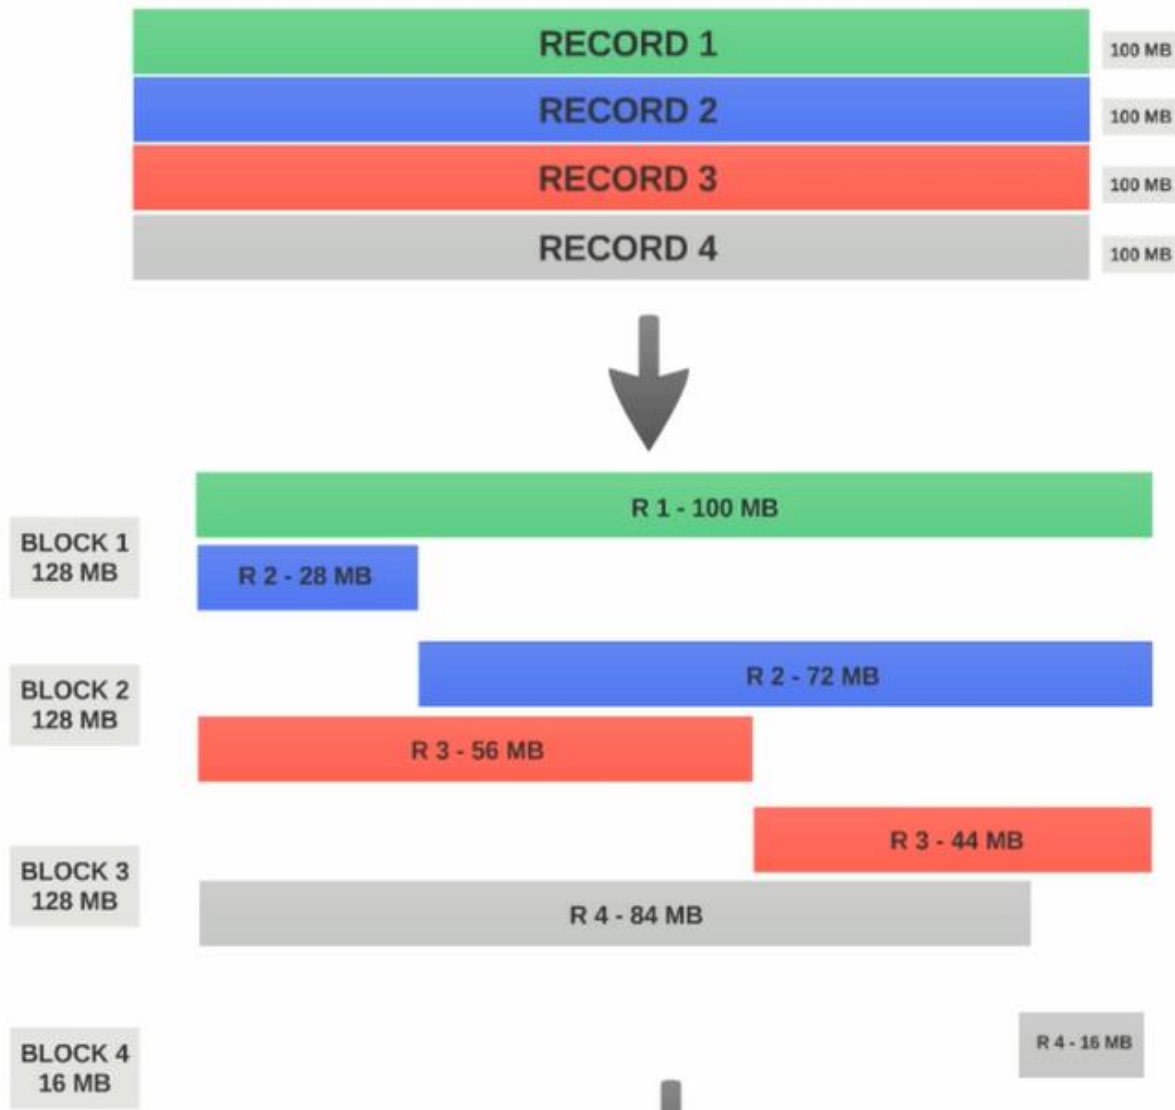
- Each mapper processes one record at a time
- Each mapper executes the same set of code on each record
- The output of the mapper will be a key-value pair

## What is the input split?

---

- Same as the block?
  - No!
- Input split is not the same as the block
- A block is a hard division of data
- So a block can end BEFORE a record ends
  - Remember a record is a single line of data (stock information for one day in our example)

# Graphically



## What is the input split?

---

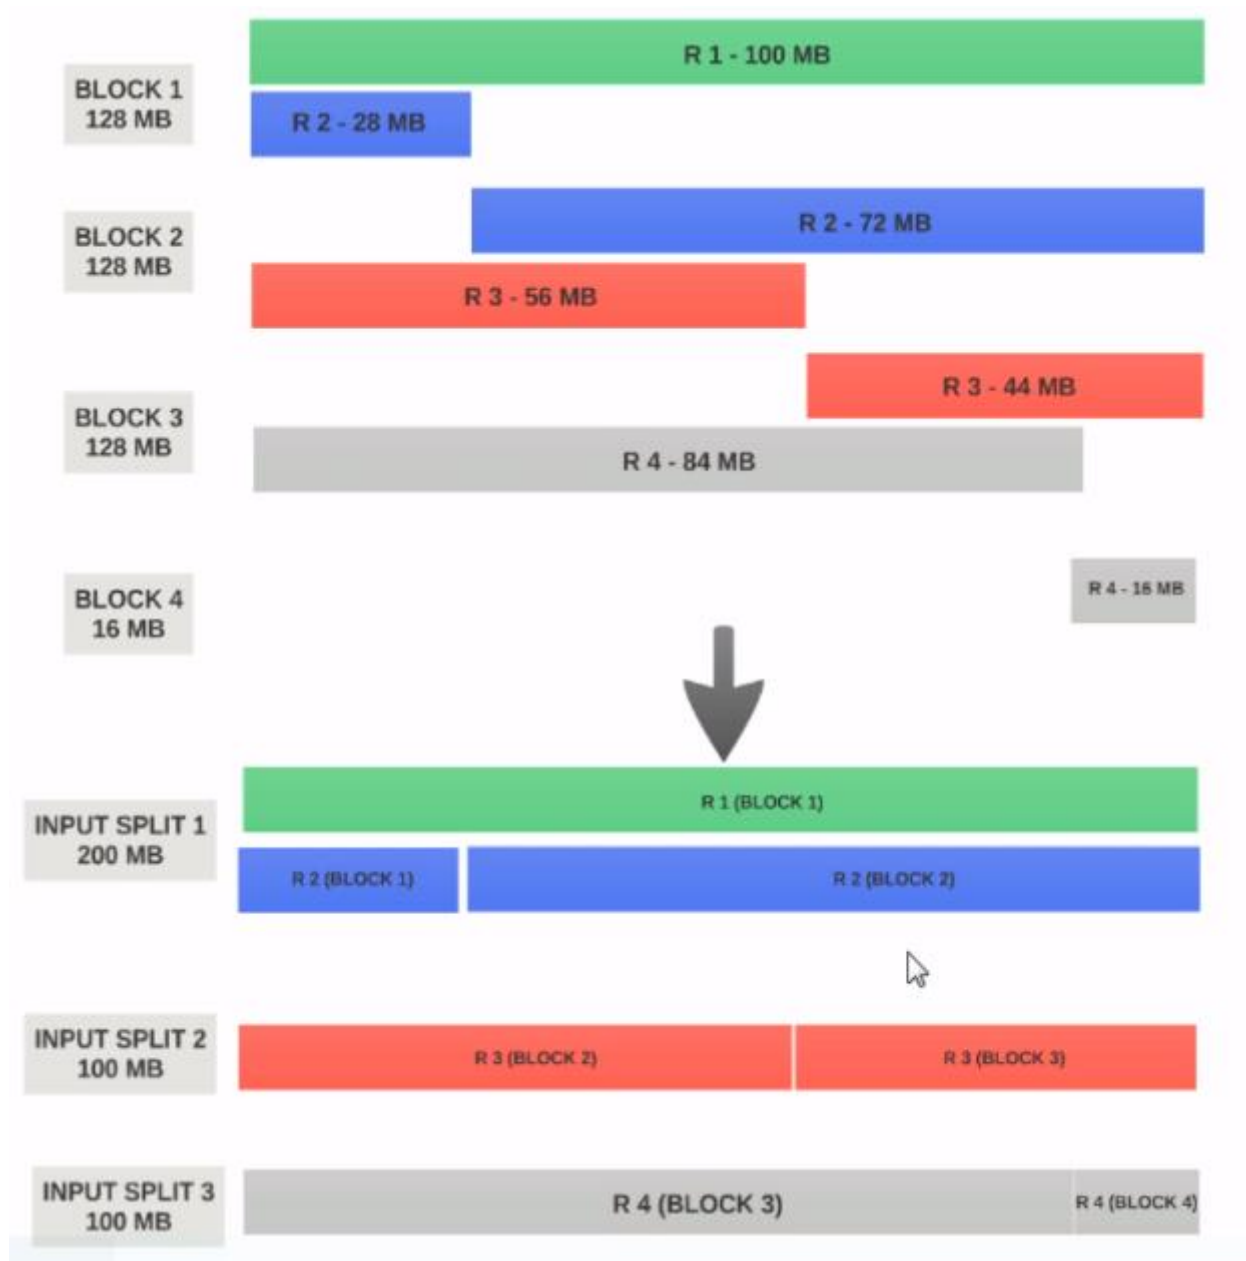
- If you assign a Mapper to Block 1, the mapper cannot process Record 2, because block 1 only has a partial record
- Input split *is not physical chunks* of data
- It is an abstraction
  - A Java class behind the scenes with pointers to start and end location within blocks
- When a mapper tries to read the data, it knows where to start and where to end



## What is the input split?

---

- The start position of a input split can start in one block and end in another block
- So this is why we have the concept of input split
- Input split *respects logical record boundaries*
- During MapReduce executions, Hadoop goes through the blocks and creates input splits that respect record boundaries



# Map Phase

---

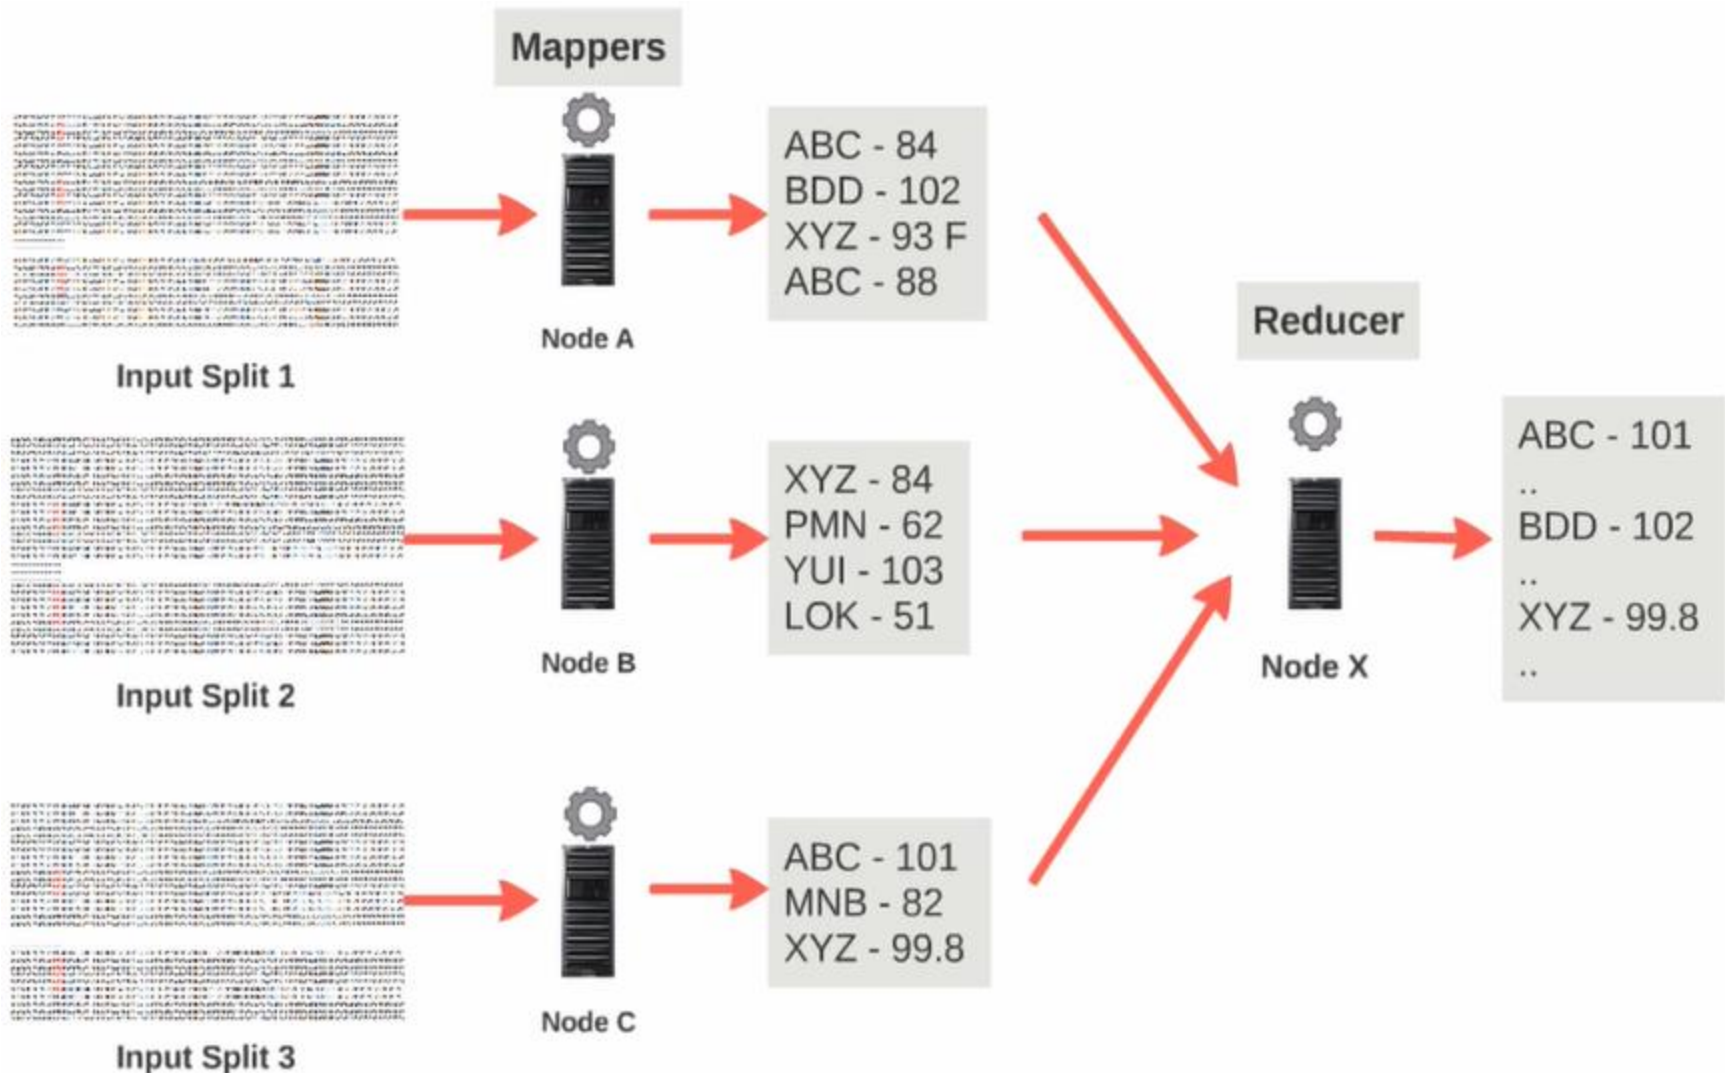
- Mapper is a program
  - Can be written in many different programming languages
  - In our case, it will be Java
- Mapper is therefore a Java program that is invoked by the Hadoop framework ***once per every record in the input split***
- So if you have 10 records in your input split, the mapper will be executed 10 times

# Interview Question

---

- How many mappers does Hadoop create for a Data Set?
  - Entirely dependent on the number of input splits
  - A mapper is invoked for every single record in the input split

# Map Phase



## Map Phase

---

- Again, a mapper is invoked for every single record in the input split
- The output of the mapper should be a key-value pair
- In our sample data set, every line is a record
- How do you decide what should be the key and the value for the key?
  - The reduce phase will give you an answer

# Reduce Phase

## Map Phase

Mapper



Node A

ABC - 60  
STT - 82  
ABC - 50  
BDD - 123  
ABC - 111  
STT - 93  
BDD - 63  
PMN - 23

Mapper



Node B

STT - 63  
ABC - 100  
BRX - 73  
ABC - 31  
STT - 115  
BRX - 103

## Reduce Phase

Reduce Function



Node X

ABC - 111  
BDD - 123  
PMN - 23  
BRX - 103  
STT - 93

ABC (60, 50, 111, 100, 31)  
BDD (123, 63)  
PMN (23)  
BRX (73, 103)  
STT (82, 93, 63, 115)

## Reduce Phase

---

- The reducers work on the output of the mappers
- The output of the individual mappers are grouped by the key and passed to the reducer
- Reducer will receive a key and list of values as input
  - The keys will be grouped



# Interview Question

---

- How many reducers is there for a given job?
  - Set by the user
- Reducer can be a bottle neck
- So it is advisable to have more than one reducer

# Shuffle

---

- What happens if you have more than one reducer?
- How does the output of the individual mappers get grouped by symbols and reach reducer?
- The magic happens in the **shuffle phase**
- This is the process in which the output of the mappers is transferred to the reducers is known as a shuffle

# Shuffle

## Map Phase

Mapper



Node A

ABC - 60  
STT - 82  
ABC - 50  
BDD - 123  
ABC - 111  
STT - 93  
BDD - 63  
PMN - 23

Mapper



Node B

STT - 63  
ABC - 100  
BRX - 73  
ABC - 31  
STT - 115  
BRX - 103

## Shuffle Phase

Sort

ABC - 60  
ABC - 50  
ABC - 111  
BDD - 123  
BDD - 63  
PMN - 23  
STT - 82  
STT - 93

Copy

ABC (60, 50, 111, 100, 31)  
PMN (23)  
BRX (73, 103)

Merge

BDD (123, 63)  
STT (82, 93, 63, 115)

## Reduce Phase

Reducer



Node X

ABC - 111  
PMN - 23  
BRX - 103

Reducer



Node Y

BDD - 123  
STT - 93

# Shuffle

---

- Each mapper will process all the assigned records in its input split and output a key value pair
- Keys in one mapper can be found in multiple mappers

## Shuffle Phase sort

---

- In the map phase, each key is assigned to a **partition** by a class called partitioner
  - One partition per mapper
- Within each partition, the key-value pairs will be **sorted by key**
- Once the key-value pairs are sorted, the key-value pairs are then copied to the appropriate reducers

# Shuffle

## Map Phase

Mapper



Node A

ABC - 60  
STT - 82  
ABC - 50  
BDD - 123  
ABC - 111  
STT - 93  
BDD - 63  
PMN - 23

Mapper



Node B

STT - 63  
ABC - 100  
BRX - 73  
ABC - 31  
STT - 115  
BRX - 103

## Shuffle Phase

Sort

Copy

Merge

ABC - 60  
ABC - 50  
ABC - 111  
BDD - 123  
BDD - 63  
PMN - 23  
STT - 82  
STT - 93

ABC - 100  
ABC - 31  
BRX - 103  
BRX - 73  
STT - 63  
STT - 115

ABC (60, 50, 111, 100, 31)  
PMN (23)  
BRX (73, 103)

BDD (123, 63)  
STT (82, 93, 63, 115)

## Reduce Phase

Reducer



Node X

ABC - 111  
PMN - 23  
BRX - 103

Reducer



Node Y

BDD - 123  
STT - 93

## Shuffle Phase copy

---

- Once the key-value pairs are sorted, the key-value pairs are then copied to the appropriate reducers

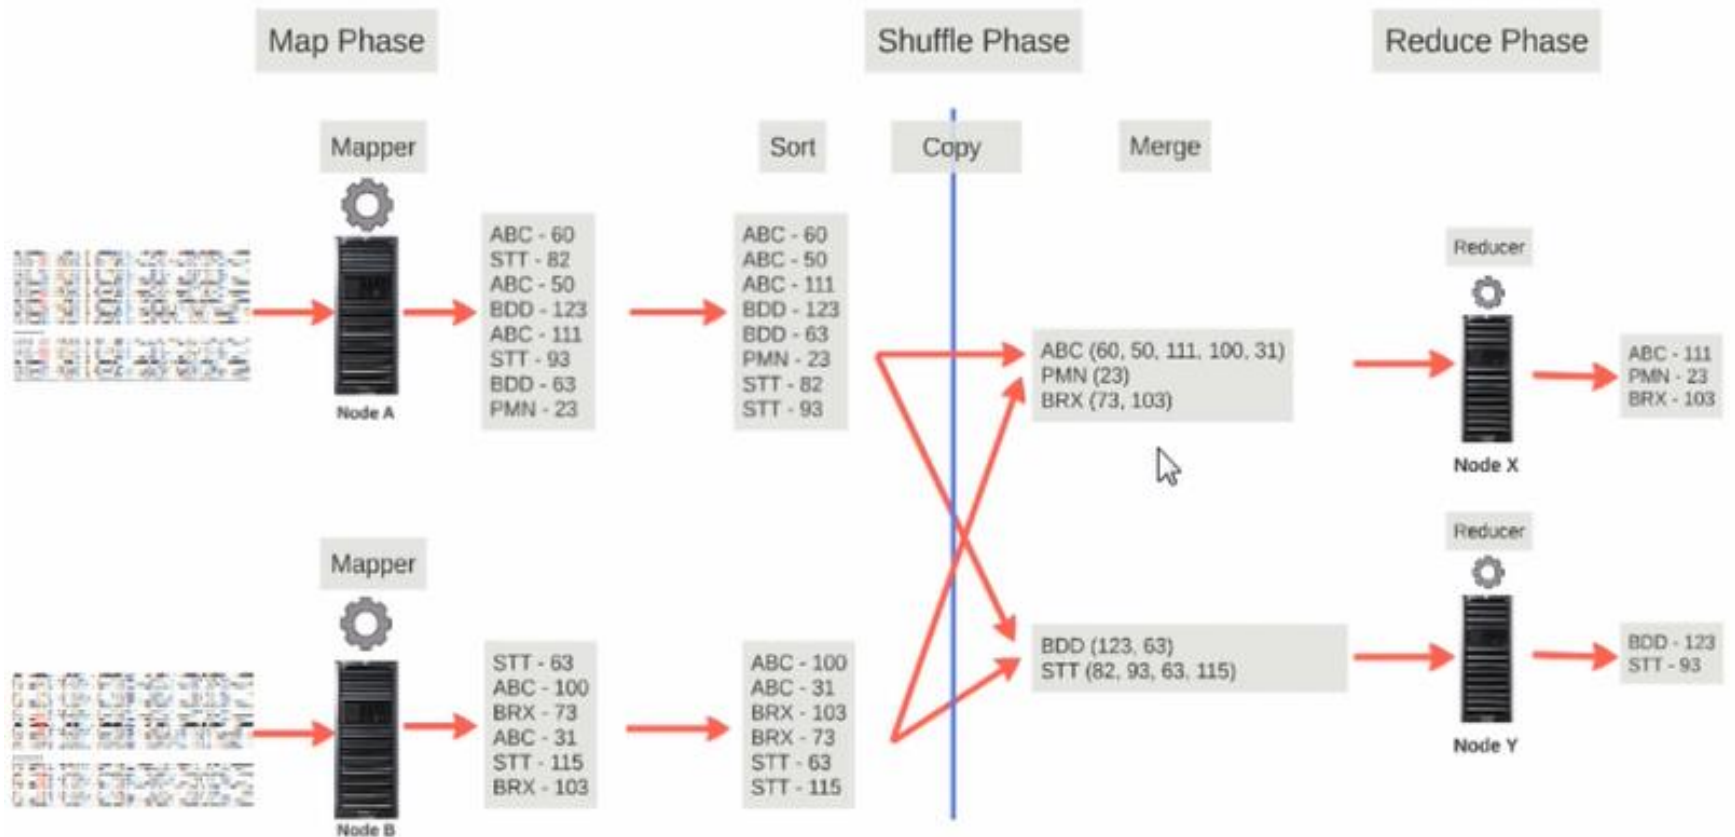
# Shuffle Merge

---

- Remember the data for partition one can come from many mappers
- The key-value pairs from different mappers will be merged maintaining the sort order
  - Hadoop framework ***guarantees*** that input the reducer is sorted by key
- Key's will be unique to each reducer



# Shuffle

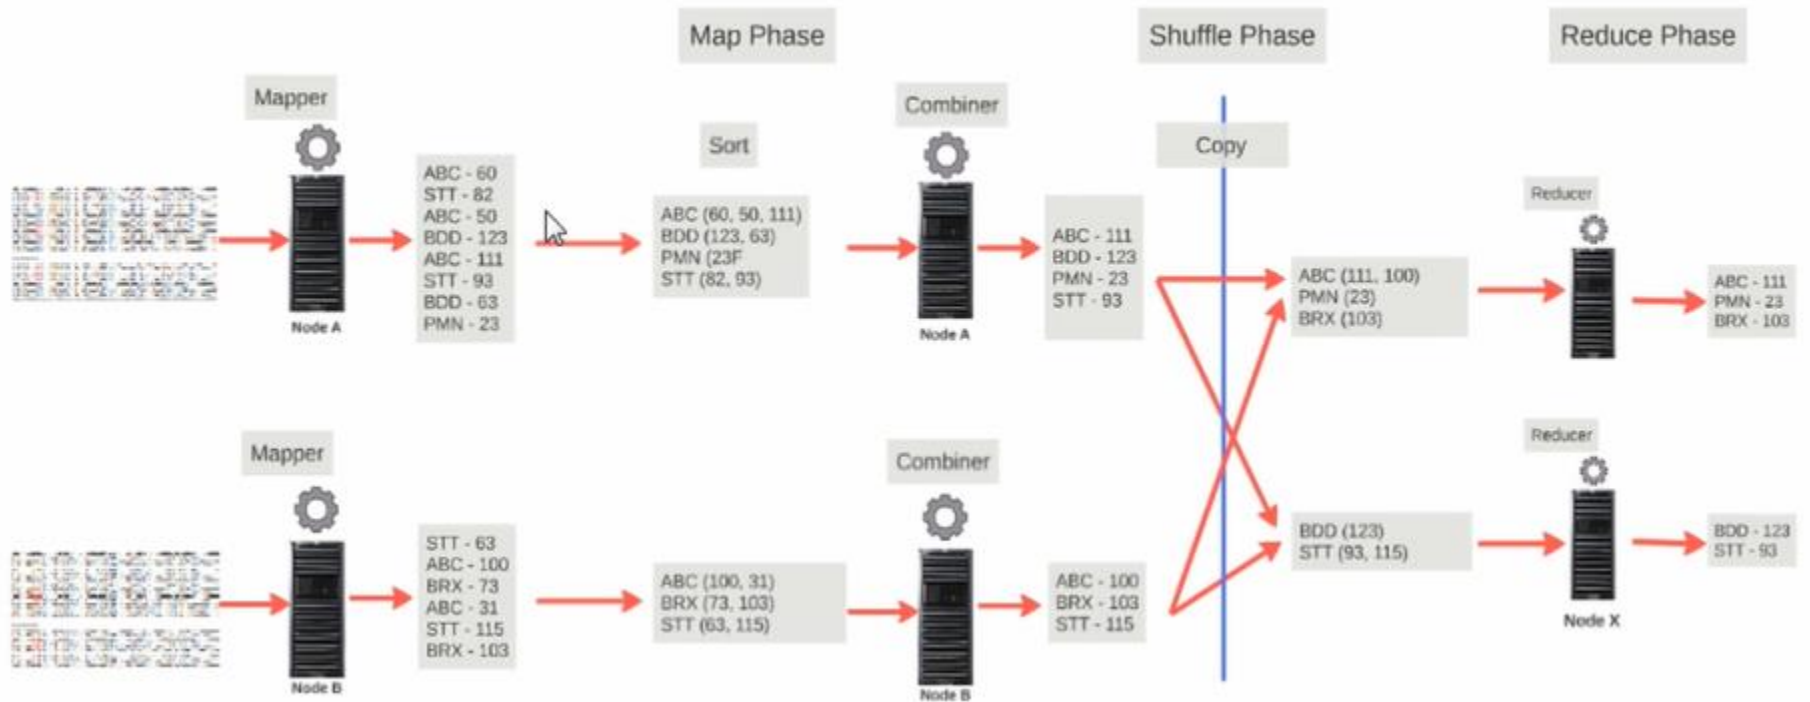


# Combiner

---

- We can also have a optional combiner at the map phase
- Combiner can be used to reduce the amount of data that is given to the reducer
- Combiner is like a mini-reducer that runs at the map phase

# Combiner



# Interview Question

---

- Can you use a combiner in all scenarios?

# MapReduce Tutorial

---

- A "Hello World" example
- Take a word file, and find the number of times words appear in the file
- Follow tutorials on:

<https://github.com/saberamini/MapReduceTutorial>